

How to Use Sklearn-genetic-opt

Introduction

Sklearn-genetic-opt uses evolutionary algorithms to fine-tune scikit-learn machine learning algorithms and perform feature selection. It is designed to accept a [scikit-learn](#) regression or classification model (or a pipeline containing one of those).

The idea behind this package is to define the set of hyperparameters we want to tune and what are their lower and upper bounds on the values they can take. It is possible to define different optimization algorithms, callbacks and build-in parameters to control how the optimization is taken. To get started, we'll use only the most basic features and options.

The optimization is made by evolutionary algorithms with the help of the [deap](#) package. It works by defining the set of hyperparameters to tune, it starts with a randomly sampled set of options (population). Then by using evolutionary operators as the mating, mutation, selection and evaluation, it generates new candidates looking to improve the cross-validation score in each generation. It'll continue with this process until a number of generations is reached or until a callback criterion is met.

Fine-tuning Example

First let's import some dataset and other scikit-learn standard modules, we'll use the [digits dataset](#). This is a classification problem, we'll fine-tune a Random Forest Classifier for this task.

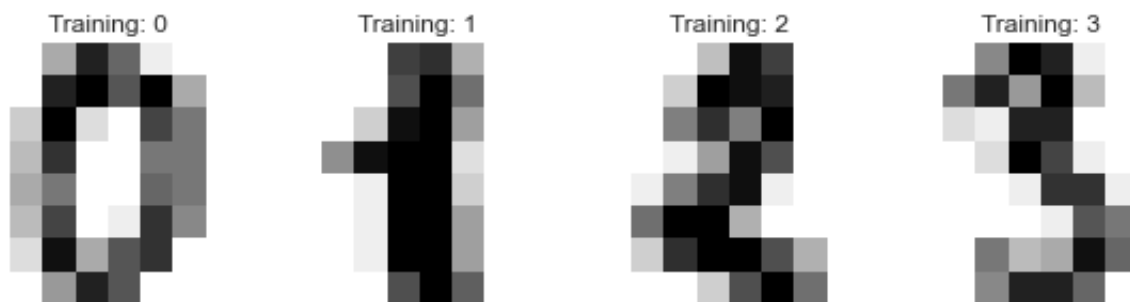
```
import matplotlib.pyplot as plt
from sklearn_genetic import GASearchCV
from sklearn_genetic.space import Categorical, Integer, Continuous
from sklearn.model_selection import train_test_split, StratifiedKFold
from sklearn.ensemble import RandomForestClassifier
from sklearn.datasets import load_digits
from sklearn.metrics import accuracy_score
```

Let's first read the data, split it in our training and test set and visualize some of the data points:

```
data = load_digits()
n_samples = len(data.images)
X = data.images.reshape((n_samples, -1))
y = data['target']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,
random_state=42)

_, axes = plt.subplots(nrows=1, ncols=4, figsize=(10, 3))
for ax, image, label in zip(axes, data.images, data.target):
    ax.set_axis_off()
    ax.imshow(image, cmap=plt.cm.gray_r, interpolation='nearest')
    ax.set_title('Training: %i' % label)
```

We should see something like this:



Now, we must define our `param_grid`, similar to scikit-learn, which is a dictionary with the model's hyperparameters. The main difference with for example scikit-learn's `GridSearchCv`, is that we don't pre-define the values to use in the search, but rather, the boundaries of each parameter.

So if we have a parameter named '`n_estimators`' we'll only tell to sklearn-genetic-opt, that is an integer value, and that we want to set a lower boundary of 100 and an upper boundary of 500, so the optimizer will set a value in this range. We must do this with all the hyperparameters we want to tune, like this:

```
param_grid = {'min_weight_fraction_leaf': Continuous(0.01, 0.5,
distribution='log-uniform'),
              'bootstrap': Categorical([True, False]),
              'max_depth': Integer(2, 30),
              'max_leaf_nodes': Integer(2, 35),
              'n_estimators': Integer(100, 300)}
```

Notice that in the case of *'bootstrap'*, as it is a categorical variable, we do must define all its possible values. As well, in the *'min_weight_fraction_leaf'*, we used an additional parameter named *distribution*, this is useful to tell the optimizer from which data distribution it can sample some random values during the optimization.

Now, we are ready to set the *GASearchCV*, its the object that will allow us to run the fitting process using evolutionary algorithms It has several options that we can use, for this first example, we'll keep it very simple:

```
# The base classifier to tune
clf = RandomForestClassifier()

# Our cross-validation strategy (it could be just an int)
cv = StratifiedKFold(n_splits=3, shuffle=True)

# The main class from sklearn-genetic-opt
evolved_estimator = GASearchCV(estimator=clf,
                              cv=cv,
                              scoring='accuracy',
                              param_grid=param_grid,
                              n_jobs=-1,
                              verbose=True)
```

So now the setup is ready, note that are other parameters that can be specified in *GASearchCV*, the ones we used, are equivalents to the meaning in scikit-learn, besides the one already explained, is worth mentioning that the “metric” is going to be used as the optimization variable, so the algorithm will try to find the set of parameters that maximizes this metric.

We are ready to run the optimization routine:

```
# Train and optimize the estimator
evolved_estimator.fit(X_train, y_train)
```

During the training process, you should see a log like this:

1 evolved_estimator.fit(X_train,y_train)					
gen	nevals	fitness	fitness_std	fitness_max	fitness_min
0	10	0.799258	0.123352	0.920143	0.541734
1	14	0.869397	0.0468278	0.909021	0.799679
2	20	0.902797	0.0214169	0.922088	0.862663
3	17	0.923556	0.00421409	0.929485	0.918343
4	19	0.923001	0.00551673	0.929485	0.910946
5	17	0.926705	0.00432808	0.929485	0.918343
6	19	0.925591	0.00486938	0.929485	0.918343
7	19	0.926145	0.00516765	0.929485	0.916511
8	19	0.923742	0.00628633	0.929485	0.914618
9	19	0.921704	0.00592012	0.929485	0.910935
10	16	0.925788	0.00405725	0.929485	0.918343
11	19	0.922245	0.0077258	0.929485	0.905328
12	19	0.918746	0.00703819	0.931357	0.909114
13	19	0.924705	0.00123214	0.927664	0.92396
14	20	0.924309	0.00262313	0.927664	0.920174
15	17	0.925981	0.00410335	0.931368	0.918405
16	17	0.928951	0.00462616	0.938806	0.920226
17	17	0.935635	0.00343967	0.938827	0.929495
18	17	0.938045	0.00170187	0.940648	0.935051
19	20	0.939145	0.00258897	0.946182	0.936923
20	18	0.941001	0.0048046	0.946182	0.931409
21	17	0.94137	0.00439765	0.948045	0.936872
22	18	0.941565	0.00663459	0.949897	0.929526
23	18	0.94618	0.0046269	0.949897	0.938775
24	19	0.945076	0.00525547	0.949897	0.936954
25	17	0.940621	0.00974303	0.949897	0.927623
26	16	0.944328	0.00563166	0.949897	0.936923
27	16	0.945627	0.00538466	0.949897	0.936923
28	17	0.944888	0.00615387	0.949897	0.936913
29	18	0.94378	0.00682917	0.949897	0.931419
30	18	0.943032	0.008666	0.949897	0.927612
31	18	0.945633	0.00658267	0.949897	0.933209
32	17	0.943215	0.00704567	0.949897	0.931337
33	14	0.943595	0.0062123	0.949897	0.933209
34	18	0.942837	0.00591748	0.949897	0.93504
35	20	0.94378	0.00579118	0.949897	0.935133

This log, shows us the metrics obtained in each iteration (generation), this is what each entry means:

- **gen:** The number of the generation
- **nevals:** How many hyperparameters were fitted in this generation
- **fitness:** The average score metric in the cross-validation (validation set). In this case, the average accuracy across the folds of all the hyperparameters sets.
- **fitness_std:** The standard deviation of the cross-validations accuracy.
- **fitness_max:** The maximum individual score of all the models in this generation.
- **fitness_min:** The minimum individual score of all the models in this generation.

After fitting the model, we have some extra methods to use the model right away. It will use by default the best set of hyperparameters it found, based on the cross-validation score:

```
# Best parameters found
print(evolved_estimator.best_params_)
# Use the model fitted with the best parameters
y_predict_ga = evolved_estimator.predict(X_test)
print(accuracy_score(y_test, y_predict_ga))
```

In this case, we got an accuracy score in the test set of 0.93

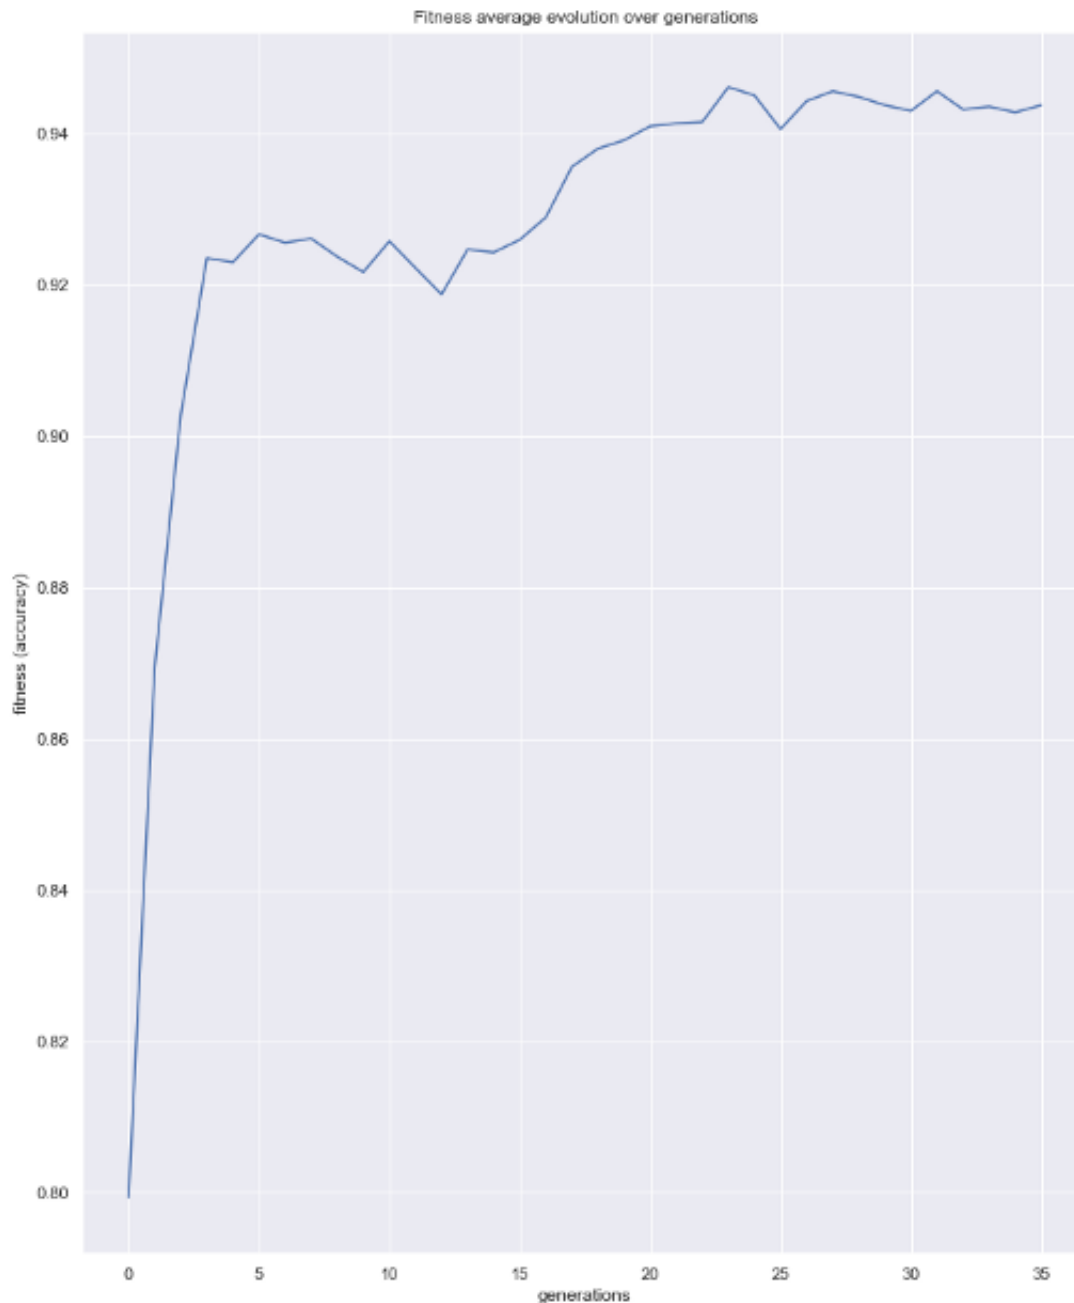
```
1 y_predict_ga = evolved_estimator.predict(X_test)
2 accuracy_score(y_test, y_predict_ga)
```

```
0.9340222575516693
```

```
1 evolved_estimator.best_params_
{'min_weight_fraction_leaf': 0.014725004803419667,
 'bootstrap': True,
 'max_depth': 25,
 'max_leaf_nodes': 29,
 'n_estimators': 259}
```

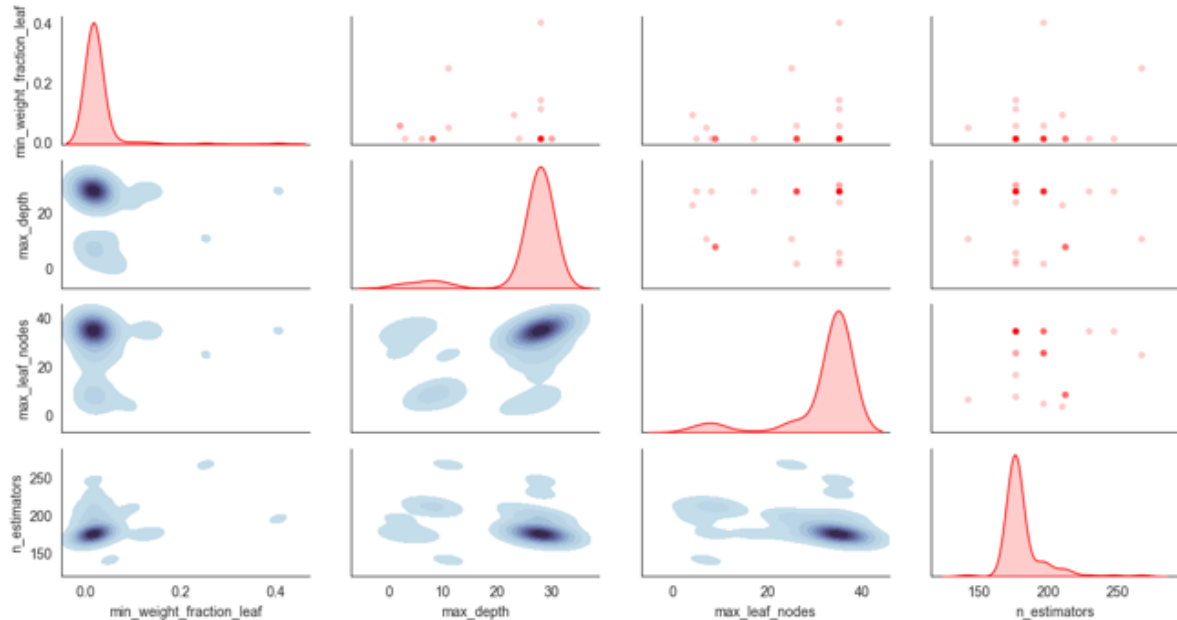
Now, let's use a couple more functions available in the package. The first one will help us to see the evolution of our metric over the generations

```
from sklearn_genetic.plots import plot_fitness_evolution
plot_fitness_evolution(evolved_estimator)
plt.show()
```



At last, we can check the property called `evolved_estimator.logbook`, this is a DEAP's logbook which stores all the results of every individual fitted model. sklearn-genetic-opt comes with a plot function to analyze this log:

```
from sklearn_genetic.plots import plot_search_space
plot_search_space(evolved_estimator, features=['min_weight_fraction_leaf',
'max_depth', 'max_leaf_nodes', 'n_estimators'])
plt.show()
```



What this plot shows us, is the distribution of the sampled values for each hyperparameter. We can see for example in the `'min_weight_fraction_leaf'` that the algorithm mostly sampled values below 0.15. You can also check every single combination of variables and the contour plot that represents the sampled values.

Feature Selection Example

For this example, we are going to use the well-known Iris dataset, it's a classification problem with four features. We are also going to simulate some random noise to represent non-important features:

```
import matplotlib.pyplot as plt
from sklearn_genetic import GAFeatureSelectionCV
from sklearn_genetic.plots import plot_fitness_evolution
from sklearn.model_selection import train_test_split, StratifiedKFold
from sklearn.svm import SVC
from sklearn.datasets import load_iris
from sklearn.metrics import accuracy_score
import numpy as np

data = load_iris()
X, y = data["data"], data["target"]

noise = np.random.uniform(0, 10, size=(X.shape[0], 10))

X = np.hstack((X, noise))

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,
                                                    random_state=0)
```

This should give us 10 extra noisy features with our train and test set.

Now we can create the GAFeatureSelectionCV object, it's very similar to the GASearchCV and they share most of the parameters, the main difference is GAFeatureSelectionCV doesn't run hyperparameters optimization thus the param_grid parameter it's not available, and the estimator should be defined with its hyperparameters.

The way the feature selection is performed is by creating models with a subsample of features and evaluate its cv-score, the way the subsets are created is by using the available evolutionary algorithms. It also tries to minimize the number of selected features, so it's a multi-objective optimization.

Let's create the feature selection object, the estimator we're going to use is a SVM:

```
clf = SVC(gamma='auto')

evolved_estimator = GAFeatureSelectionCV(
    estimator=clf,
    cv=3,
    scoring="accuracy",
    population_size=30,
    generations=20,
    n_jobs=-1,
    verbose=True,
    keep_top_k=2,
    elitism=True,
)
```

We are ready to run the optimization routine:

```
# Train and select the features
evolved_estimator.fit(X_train, y_train)
```

During the training, the same log format is displayed as before:

gen	nevals	fitness	fitness_std	fitness_max	fitness_min
0	30	0.558444	0.155441	0.893333	0.253333
1	54	0.659333	0.132948	0.893333	0.333333
2	54	0.742667	0.0867111	0.893333	0.586667
3	55	0.805778	0.0740117	0.893333	0.653333
4	52	0.873333	0.0435125	0.906667	0.746667
5	53	0.896222	0.00659592	0.913333	0.893333
6	55	0.901111	0.0131186	0.953333	0.893333
7	54	0.911778	0.0206332	0.953333	0.893333
8	50	0.926444	0.0210455	0.953333	0.893333
9	51	0.941333	0.020177	0.966667	0.913333
10	49	0.955556	0.00978787	0.966667	0.913333
11	55	0.959111	0.00660714	0.966667	0.953333
12	57	0.965333	0.004	0.966667	0.953333
13	55	0.966444	0.00271257	0.973333	0.953333
14	58	0.966667	6.66134e-16	0.966667	0.966667
15	53	0.966889	0.0011967	0.973333	0.966667
16	56	0.967556	0.00226623	0.973333	0.966667
17	53	0.969556	0.00330357	0.973333	0.966667
18	51	0.971111	0.0031427	0.973333	0.966667
19	58	0.972889	0.00166296	0.973333	0.966667
20	54	0.973333	3.33067e-16	0.973333	0.973333

After fitting the model, we have some extra methods to use the model right away. It will use by default the best set of features it found, remember as the algorithm used only a subset, you have to select them from the `X_test array`, this is done like this:

```
features = evolved_estimator.best_features_

# Predict only with the subset of selected features
y_predict_ga = evolved_estimator.predict(X_test[:, features])
accuracy = accuracy_score(y_test, y_predict_ga)
```

```
print(evolved_estimator.best_features_)
print("accuracy score: ", "{:.2f}".format(accuracy))

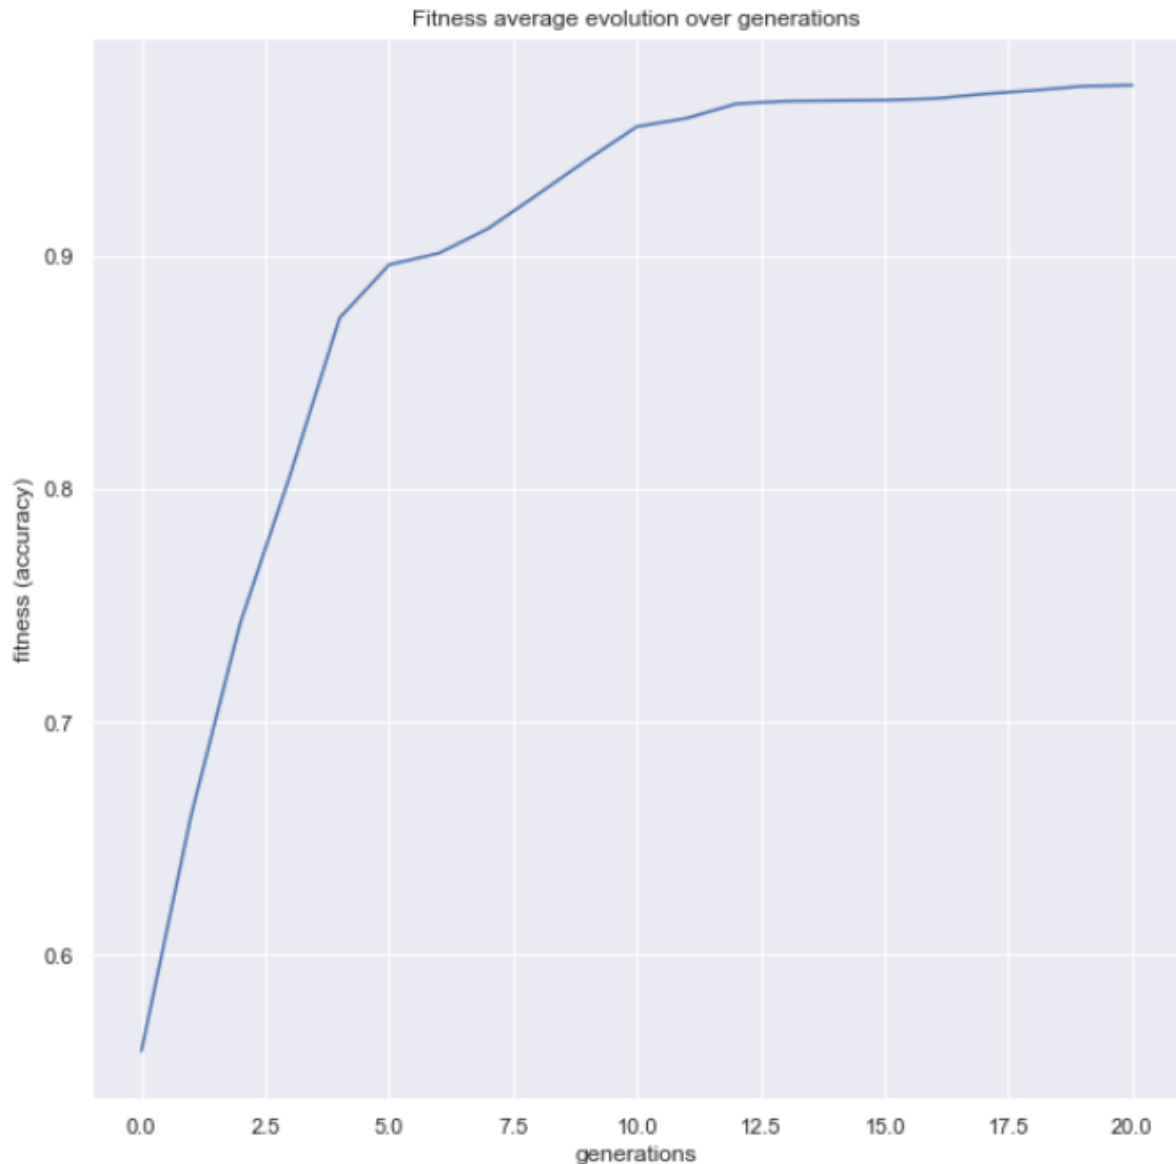
[ True  True  True  True False False False False False False False
  False False]
accuracy score:  0.98
```

In this case, we got an accuracy score in the test set of 0.98.

Notice that the `best_features_` is a vector of bool values, each position represents the index of the feature (column) and the value indicates if that features was selected (True) or not (False) by the algorithm. In this example, the algorithm, discarded all the noisy random variables we created and selected the original variables.

We can also plot the fitness evolution:

```
from sklearn_genetic.plots import plot_fitness_evolution
plot_fitness_evolution(evolved_estimator)
plt.show()
```



This concludes our introduction to the basic sklearn-genetic-opt usage. Further tutorials will cover the `GASearchCV` and `GAFeatureSelectionCV` parameters, callbacks, different optimization algorithms and more advanced use cases.