

sklearn.linear_model.LogisticRegression

```
class sklearn.linear_model.LogisticRegression(penalty='l2', *, dual=False, tol=0.0001, C=1.0, fit_intercept=True,
intercept_scaling=1, class_weight=None, random_state=None, solver='lbfgs', max_iter=100, multi_class='auto', verbose=0,
warm_start=False, n_jobs=None, l1_ratio=None)
```

[\[source\]](#)

Logistic Regression (aka logit, MaxEnt) classifier.

In the multiclass case, the training algorithm uses the one-vs-rest (OvR) scheme if the 'multi_class' option is set to 'ovr', and uses the cross-entropy loss if the 'multi_class' option is set to 'multinomial'. (Currently the 'multinomial' option is supported only by the 'lbfgs', 'sag', 'saga' and 'newton-cg' solvers.)

This class implements regularized logistic regression using the 'liblinear' library, 'newton-cg', 'sag', 'saga' and 'lbfgs' solvers. **Note that regularization is applied by default.** It can handle both dense and sparse input. Use C-ordered arrays or CSR matrices containing 64-bit floats for optimal performance; any other input format will be converted (and copied).

The 'newton-cg', 'sag', and 'lbfgs' solvers support only L2 regularization with primal formulation, or no regularization. The 'liblinear' solver supports both L1 and L2 regularization, with a dual formulation only for the L2 penalty. The Elastic-Net regularization is only supported by the 'saga' solver.

Read more in the [User Guide](#).

Parameters:

penalty : {'l1', 'l2', 'elasticnet', 'none'}, default='l2'

Specify the norm of the penalty:

- 'none': no penalty is added;
- 'l2': add a L2 penalty term and it is the default choice;
- 'l1': add a L1 penalty term;
- 'elasticnet': both L1 and L2 penalty terms are added.

Warning: Some penalties may not work with some solvers. See the parameter `solver` below, to know the compatibility between the penalty and solver.

New in version 0.19: l1 penalty with SAGA solver (allowing 'multinomial' + L1)

dual : bool, default=False

Dual or primal formulation. Dual formulation is only implemented for l2 penalty with liblinear solver. Prefer dual=False when `n_samples > n_features`.

tol : float, default=1e-4

Tolerance for stopping criteria.

C : float, default=1.0

Inverse of regularization strength; must be a positive float. Like in support vector machines, smaller values specify stronger regularization.

fit_intercept : bool, default=True

Specifies if a constant (a.k.a. bias or intercept) should be added to the decision function.

intercept_scaling : float, default=1

Useful only when the solver 'liblinear' is used and `self.fit_intercept` is set to True. In this case, `x` becomes `[x, self.intercept_scaling]`, i.e. a "synthetic" feature with constant value equal to `intercept_scaling` is appended to the instance vector. The intercept becomes `intercept_scaling * synthetic_feature_weight`.

Note! the synthetic feature weight is subject to l1/l2 regularization as all other features. To lessen the effect of regularization on synthetic feature weight (and therefore on the intercept) `intercept_scaling` has to be increased.

`class_weight` : *dict* or *'balanced'*, *default=None*

Weights associated with classes in the form `{class_label: weight}`. If not given, all classes are supposed to have weight one.

The “balanced” mode uses the values of `y` to automatically adjust weights inversely proportional to class frequencies in the input data as `n_samples / (n_classes * np.bincount(y))`.

Note that these weights will be multiplied with `sample_weight` (passed through the fit method) if `sample_weight` is specified.

New in version 0.17: `class_weight='balanced'`

`random_state` : *int*, *RandomState* instance, *default=None*

Used when `solver == 'sag', 'saga'` or `'liblinear'` to shuffle the data. See [Glossary](#) for details.

`solver` : *{'newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga'}*, *default='lbfgs'*

Algorithm to use in the optimization problem. Default is `'lbfgs'`. To choose a solver, you might want to consider the following aspects:

- For small datasets, `'liblinear'` is a good choice, whereas `'sag'` and `'saga'` are faster for large ones;
- For multiclass problems, only `'newton-cg'`, `'sag'`, `'saga'` and `'lbfgs'` handle multinomial loss;
- `'liblinear'` is limited to one-versus-rest schemes.

Warning: The choice of the algorithm depends on the penalty chosen: Supported penalties by solver:

- `'newton-cg'` - `['l2', 'none']`
- `'lbfgs'` - `['l2', 'none']`
- `'liblinear'` - `['l1', 'l2']`
- `'sag'` - `['l2', 'none']`
- `'saga'` - `['elasticnet', 'l1', 'l2', 'none']`

Note: `'sag'` and `'saga'` fast convergence is only guaranteed on features with approximately the same scale. You can preprocess the data with a scaler from [sklearn.preprocessing](#).

See also: Refer to the User Guide for more information regarding [LogisticRegression](#) and more specifically the [Table summarizing solver/penalty supports](#). <!-- # noqa: E501 -->

New in version 0.17: Stochastic Average Gradient descent solver.

New in version 0.19: SAGA solver.

Changed in version 0.22: The default solver changed from `'liblinear'` to `'lbfgs'` in 0.22.

`max_iter` : *int*, *default=100*

Maximum number of iterations taken for the solvers to converge.

`multi_class` : *{'auto', 'ovr', 'multinomial'}*, *default='auto'*

If the option chosen is `'ovr'`, then a binary problem is fit for each label. For `'multinomial'` the loss minimised is the multinomial loss fit across the entire probability distribution, *even when the data is binary*. `'multinomial'` is unavailable when `solver='liblinear'`. `'auto'` selects `'ovr'` if the data is binary, or if `solver='liblinear'`, and otherwise selects `'multinomial'`.

New in version 0.18: Stochastic Average Gradient descent solver for `'multinomial'` case.

Changed in version 0.22: Default changed from `'ovr'` to `'auto'` in 0.22.

`verbose` : *int*, *default=0*

For the liblinear and lbfgs solvers set verbose to any positive number for verbosity.

warm_start : bool, default=False

When set to True, reuse the solution of the previous call to fit as initialization, otherwise, just erase the previous solution. Useless for liblinear solver. See [the Glossary](#).

New in version 0.17: warm_start to support lbfgs, newton-cg, sag, saga solvers.

n_jobs : int, default=None

Number of CPU cores used when parallelizing over classes if multi_class='ovr'. This parameter is ignored when the solver is set to 'liblinear' regardless of whether 'multi_class' is specified or not. None means 1 unless in a [joblib.parallel_backend](#) context. -1 means using all processors. See [Glossary](#) for more details.

l1_ratio : float, default=None

The Elastic-Net mixing parameter, with $0 \leq \text{l1_ratio} \leq 1$. Only used if penalty='elasticnet'. Setting l1_ratio=0 is equivalent to using penalty='l2', while setting l1_ratio=1 is equivalent to using penalty='l1'. For $0 < \text{l1_ratio} < 1$, the penalty is a combination of L1 and L2.

Attributes:

classes_ : ndarray of shape (n_classes,)

A list of class labels known to the classifier.

coef_ : ndarray of shape (1, n_features) or (n_classes, n_features)

Coefficient of the features in the decision function.

coef_ is of shape (1, n_features) when the given problem is binary. In particular, when multi_class='multinomial', coef_ corresponds to outcome 1 (True) and -coef_ corresponds to outcome 0 (False).

intercept_ : ndarray of shape (1,) or (n_classes,)

Intercept (a.k.a. bias) added to the decision function.

If fit_intercept is set to False, the intercept is set to zero. intercept_ is of shape (1,) when the given problem is binary. In particular, when multi_class='multinomial', intercept_ corresponds to outcome 1 (True) and -intercept_ corresponds to outcome 0 (False).

n_features_in_ : int

Number of features seen during [fit](#).

New in version 0.24.

feature_names_in_ : ndarray of shape (n_features_in_,)

Names of features seen during [fit](#). Defined only when X has feature names that are all strings.

New in version 1.0.

n_iter_ : ndarray of shape (n_classes,) or (1,)

Actual number of iterations for all classes. If binary or multinomial, it returns only 1 element. For liblinear solver, only the maximum number of iteration across all classes is given.

Changed in version 0.20: In SciPy $\leq 1.0.0$ the number of lbfgs iterations may exceed max_iter. n_iter_ will now report at most max_iter.

See also:

[SGDClassifier](#)

Incrementally trained logistic regression (when given the parameter loss="log").

[LogisticRegressionCV](#)

Logistic regression with built-in cross validation.

Notes

The underlying C implementation uses a random number generator to select features when fitting the model. It is thus not uncommon, to have slightly different results for the same input data. If that happens, try with a smaller `tol` parameter.

Predict output may not match that of standalone liblinear in certain cases. See [differences from liblinear](#) in the narrative documentation.

References

L-BFGS-B – Software for Large-scale Bound-constrained Optimization

Ciyou Zhu, Richard Byrd, Jorge Nocedal and Jose Luis Morales. <http://users.iems.northwestern.edu/~nocedal/lbfgsb.html>

LIBLINEAR – A Library for Large Linear Classification

<https://www.csie.ntu.edu.tw/~cjlin/liblinear/>

SAG – Mark Schmidt, Nicolas Le Roux, and Francis Bach

Minimizing Finite Sums with the Stochastic Average Gradient <https://hal.inria.fr/hal-00860051/document>

SAGA – Defazio, A., Bach F. & Lacoste-Julien S. (2014).

[“SAGA: A Fast Incremental Gradient Method With Support for Non-Strongly Convex Composite Objectives”](#)

Hsiang-Fu Yu, Fang-Lan Huang, Chih-Jen Lin (2011). Dual coordinate descent

methods for logistic regression and maximum entropy models. Machine Learning 85(1-2):41-75.

https://www.csie.ntu.edu.tw/~cjlin/papers/maxent_dual.pdf

Examples

```
>>> from sklearn.datasets import load_iris
>>> from sklearn.linear_model import LogisticRegression
>>> X, y = load_iris(return_X_y=True)
>>> clf = LogisticRegression(random_state=0).fit(X, y)
>>> clf.predict(X[:2, :])
array([0, 0])
>>> clf.predict_proba(X[:2, :])
array([[9.8...e-01, 1.8...e-02, 1.4...e-08],
       [9.7...e-01, 2.8...e-02, ...e-08]])
>>> clf.score(X, y)
0.97...
```

Methods

<code>decision_function(X)</code>	Predict confidence scores for samples.
<code>densify()</code>	Convert coefficient matrix to dense array format.
<code>fit(X, y[, sample_weight])</code>	Fit the model according to the given training data.
<code>get_params([deep])</code>	Get parameters for this estimator.
<code>predict(X)</code>	Predict class labels for samples in X.
<code>predict_log_proba(X)</code>	Predict logarithm of probability estimates.
<code>predict_proba(X)</code>	Probability estimates.
<code>score(X, y[, sample_weight])</code>	Return the mean accuracy on the given test data and labels.
<code>set_params(**params)</code>	Set the parameters of this estimator.
<code>sparsify()</code>	Convert coefficient matrix to sparse format.

`decision_function(X)`

[\[source\]](#)

Predict confidence scores for samples.

The confidence score for a sample is proportional to the signed distance of that sample to the hyperplane.

Parameters:***X : {array-like, sparse matrix} of shape (n_samples, n_features)***

The data matrix for which we want to get the confidence scores.

Returns:***scores : ndarray of shape (n_samples,) or (n_samples, n_classes)***Confidence scores per (n_samples, n_classes) combination. In the binary case, confidence score for `self.classes_[1]` where `>0` means this class would be predicted.**`densify()`**[\[source\]](#)

Convert coefficient matrix to dense array format.

Converts the `coef_` member (back) to a `numpy.ndarray`. This is the default format of `coef_` and is required for fitting, so calling this method is only required on models that have previously been sparsified; otherwise, it is a no-op.**Returns:****`self`**

Fitted estimator.

`fit(X, y, sample_weight=None)`[\[source\]](#)

Fit the model according to the given training data.

Parameters:***X : {array-like, sparse matrix} of shape (n_samples, n_features)***Training vector, where `n_samples` is the number of samples and `n_features` is the number of features.***y : array-like of shape (n_samples,)***

Target vector relative to X.

sample_weight : array-like of shape (n_samples,) default=None

Array of weights that are assigned to individual samples. If not provided, then each sample is given unit weight.

*New in version 0.17: sample_weight support to LogisticRegression.***Returns:****`self`**

Fitted estimator.

Notes

The SAGA solver supports both float64 and float32 bit arrays.

`get_params(deep=True)`[\[source\]](#)

Get parameters for this estimator.

Parameters:**deep : bool, default=True**

If True, will return the parameters for this estimator and contained subobjects that are estimators.

Returns:**params : dict**

Parameter names mapped to their values.

predict(X)[\[source\]](#)

Predict class labels for samples in X.

Parameters:**X : {array-like, sparse matrix} of shape (n_samples, n_features)**

The data matrix for which we want to get the predictions.

Returns:**y_pred : ndarray of shape (n_samples,)**

Vector containing the class labels for each sample.

predict_log_proba(X)[\[source\]](#)

Predict logarithm of probability estimates.

The returned estimates for all classes are ordered by the label of classes.

Parameters:**X : array-like of shape (n_samples, n_features)**

Vector to be scored, where `n_samples` is the number of samples and `n_features` is the number of features.

Returns:**T : array-like of shape (n_samples, n_classes)**

Returns the log-probability of the sample for each class in the model, where classes are ordered as they are in `self.classes_`.

predict_proba(X)[\[source\]](#)

Probability estimates.

The returned estimates for all classes are ordered by the label of classes.

For a multi_class problem, if `multi_class` is set to be "multinomial" the softmax function is used to find the predicted probability of each class. Else use a one-vs-rest approach, i.e calculate the probability of each class assuming it to be positive using the logistic function. and normalize these values across all the classes.

Parameters:**X : array-like of shape (n_samples, n_features)**Vector to be scored, where `n_samples` is the number of samples and `n_features` is the number of features.**Returns:****T : array-like of shape (n_samples, n_classes)**Returns the probability of the sample for each class in the model, where classes are ordered as they are in `self.classes_`.`score(X, y, sample_weight=None)`[\[source\]](#)

Return the mean accuracy on the given test data and labels.

In multi-label classification, this is the subset accuracy which is a harsh metric since you require for each sample that each label set be correctly predicted.

Parameters:**X : array-like of shape (n_samples, n_features)**

Test samples.

y : array-like of shape (n_samples,) or (n_samples, n_outputs)

True labels for X.

sample_weight : array-like of shape (n_samples,), default=None

Sample weights.

Returns:**score : float**Mean accuracy of `self.predict(X)` wrt. `y`.`set_params(**params)`[\[source\]](#)

Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects (such as [Pipeline](#)). The latter have parameters of the form `<component>__<parameter>` so that it's possible to update each component of a nested object.**Parameters:******params : dict**

Estimator parameters.

Returns:**self : estimator instance**

Estimator instance.

`sparsify()`[\[source\]](#)

Convert coefficient matrix to sparse format.

Converts the `coef_` member to a `scipy.sparse` matrix, which for L1-regularized models can be much more memory- and storage-efficient than the usual `numpy.ndarray` representation.The `intercept_` member is not converted.

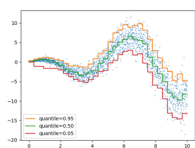
Returns:**self**

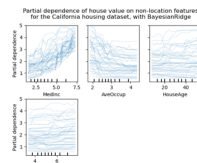
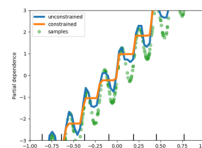
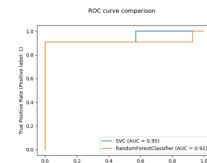
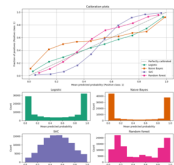
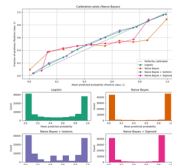
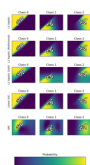
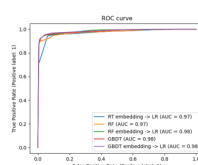
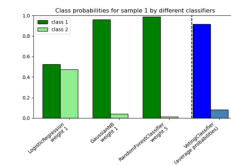
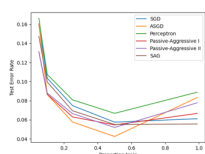
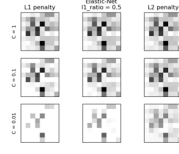
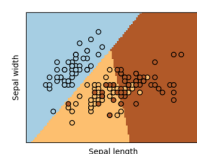
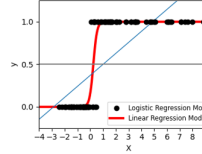
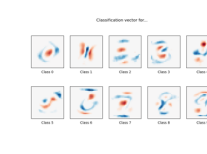
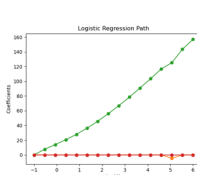
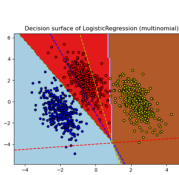
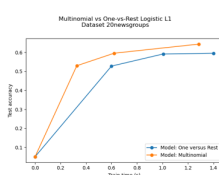
Fitted estimator.

Notes

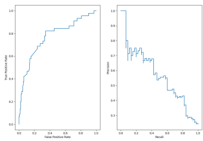
For non-sparse models, i.e. when there are not many zeros in `coef_`, this may actually *increase* memory usage, so use this method with care. A rule of thumb is that the number of zero elements, which can be computed with `(coef_ == 0).sum()`, must be more than 50% for this to provide significant benefits.

After calling this method, further fitting with the `partial_fit` method (if any) will not work until you call `densify`.

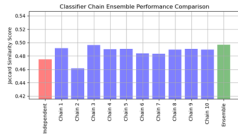
Examples using `sklearn.linear_model.LogisticRegression`
[Release Highlights for scikit-learn 1.1](#)

[Release Highlights for scikit-learn 1.0](#)

[Release Highlights for scikit-learn 0.24](#)

[Release Highlights for scikit-learn 0.23](#)

[Release Highlights for scikit-learn 0.22](#)

[Comparison of Calibration of Classifiers](#)

[Probability Calibration curves](#)

[Plot classification probability](#)

[Feature transformations with ensembles of trees](#)

[Plot class probabilities calculated by the VotingClassifier](#)

[Comparing various online solvers](#)

[L1 Penalty and Sparsity in Logistic Regression](#)

[Logistic Regression 3-class Classifier](#)

[Logistic function](#)

[MNIST classification using multinomial logistic + L1](#)


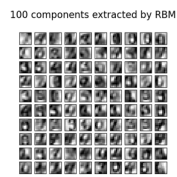
[Multiclass sparse
logistic regression
on 20newgroups](#)



[Plot multinomial and
One-vs-Rest
Logistic Regression](#)



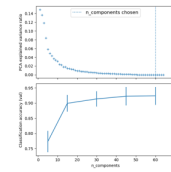
[Regularization path
of L1- Logistic
Regression](#)



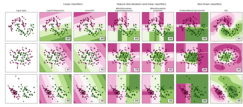
[Compact estimator
representations](#)



[Displaying Pipelines](#)



[Visualizations with
Display Objects](#)



[Classifier Chain](#)



[Restricted
Boltzmann Machine
features for digit
classification](#)

[Column Transformer
with Mixed Types](#)

[Pipelining: chaining a
PCA and a logistic
regression](#)

[Feature
discretization](#)

[Digits Classification
Exercise](#)

Toggle Menu

© 2007 - 2022, scikit-learn developers (BSD License). [Show this page source](#)