

sklearn.preprocessing.StandardScaler

```
class sklearn.preprocessing.StandardScaler(*, copy=True, with_mean=True, with_std=True)
```

[\[source\]](#)

Standardize features by removing the mean and scaling to unit variance.

The standard score of a sample x is calculated as:

$$z = (x - u) / s$$

where u is the mean of the training samples or zero if `with_mean=False`, and s is the standard deviation of the training samples or one if `with_std=False`.

Centering and scaling happen independently on each feature by computing the relevant statistics on the samples in the training set. Mean and standard deviation are then stored to be used on later data using [transform](#).

Standardization of a dataset is a common requirement for many machine learning estimators: they might behave badly if the individual features do not more or less look like standard normally distributed data (e.g. Gaussian with 0 mean and unit variance).

For instance many elements used in the objective function of a learning algorithm (such as the RBF kernel of Support Vector Machines or the L1 and L2 regularizers of linear models) assume that all features are centered around 0 and have variance in the same order. If a feature has a variance that is orders of magnitude larger than others, it might dominate the objective function and make the estimator unable to learn from other features correctly as expected.

This scaler can also be applied to sparse CSR or CSC matrices by passing `with_mean=False` to avoid breaking the sparsity structure of the data.

Read more in the [User Guide](#).

Parameters:

copy : bool, default=True

If False, try to avoid a copy and do inplace scaling instead. This is not guaranteed to always work inplace; e.g. if the data is not a NumPy array or scipy.sparse CSR matrix, a copy may still be returned.

with_mean : bool, default=True

If True, center the data before scaling. This does not work (and will raise an exception) when attempted on sparse matrices, because centering them entails building a dense matrix which in common use cases is likely to be too large to fit in memory.

with_std : bool, default=True

If True, scale the data to unit variance (or equivalently, unit standard deviation).

Attributes:

scale_ : ndarray of shape (n_features,) or None

Per feature relative scaling of the data to achieve zero mean and unit variance. Generally this is calculated using `np.sqrt(var_)`. If a variance is zero, we can't achieve unit variance, and the data is left as-is, giving a scaling factor of 1. `scale_` is equal to None when `with_std=False`.

New in version 0.17: scale_

mean_ : ndarray of shape (n_features,) or None

The mean value for each feature in the training set. Equal to None when `with_mean=False`.

var_ : ndarray of shape (n_features,) or None

The variance for each feature in the training set. Used to compute `scale_`. Equal to None when `with_std=False`.

n_features_in_ : int

Number of features seen during [fit](#).

New in version 0.24.

feature_names_in_ : ndarray of shape (n_features_in_,)

Names of features seen during [fit](#). Defined only when X has feature names that are all strings.

New in version 1.0.

n_samples_seen_ : int or ndarray of shape (n_features,)

The number of samples processed by the estimator for each feature. If there are no missing samples, the `n_samples_seen` will be an integer, otherwise it will be an array of dtype int. If `sample_weights` are used it will be a float (if no missing data) or an array of dtype float that sums the weights seen so far. Will be reset on new calls to `fit`, but increments across `partial_fit` calls.

See also:[scale](#)

Equivalent function without the estimator API.

[PCA](#)

Further removes the linear correlation across features with `'whiten=True'`.

Notes

NaNs are treated as missing values: disregarded in `fit`, and maintained in `transform`.

We use a biased estimator for the standard deviation, equivalent to `numpy.std(x, ddof=0)`. Note that the choice of `ddof` is unlikely to affect model performance.

For a comparison of the different scalers, transformers, and normalizers, see [examples/preprocessing/plot_all_scaling.py](#).

Examples

```
>>> from sklearn.preprocessing import StandardScaler
>>> data = [[0, 0], [0, 0], [1, 1], [1, 1]]
>>> scaler = StandardScaler()
>>> print(scaler.fit(data))
StandardScaler()
>>> print(scaler.mean_)
[0.5 0.5]
>>> print(scaler.transform(data))
[[-1. -1.]
 [-1. -1.]
 [ 1.  1.]
 [ 1.  1.]]
>>> print(scaler.transform([[2, 2]]))
[[3. 3.]]
```

Methods

<code>fit</code> (X[, y, sample_weight])	Compute the mean and std to be used for later scaling.
<code>fit_transform</code> (X[, y])	Fit to data, then transform it.
<code>get_feature_names_out</code> ([input_features])	Get output feature names for transformation.
<code>get_params</code> ([deep])	Get parameters for this estimator.
<code>inverse_transform</code> (X[, copy])	Scale back the data to the original representation.
<code>partial_fit</code> (X[, y, sample_weight])	Online computation of mean and std on X for later scaling.
<code>set_params</code> (**params)	Set the parameters of this estimator.
<code>transform</code> (X[, copy])	Perform standardization by centering and scaling.

[`fit`](#)(X, y=None, sample_weight=None)

[\[source\]](#)

Compute the mean and std to be used for later scaling.

Parameters:

X : {array-like, sparse matrix} of shape (n_samples, n_features)

The data used to compute the mean and standard deviation used for later scaling along the features axis.

y : None

Ignored.

sample_weight : array-like of shape (n_samples,), default=None

Individual weights for each sample.

New in version 0.24: parameter *sample_weight* support to StandardScaler.

Returns:

self : object

Fitted scaler.

[`fit_transform`](#)(X, y=None, **fit_params)

[\[source\]](#)

Fit to data, then transform it.

Fits transformer to X and y with optional parameters `fit_params` and returns a transformed version of X.

Parameters:***X : array-like of shape (n_samples, n_features)***

Input samples.

y : array-like of shape (n_samples,) or (n_samples, n_outputs), default=None

Target values (None for unsupervised transformations).

*****fit_params : dict***

Additional fit parameters.

Returns:***X_new : ndarray array of shape (n_samples, n_features_new)***

Transformed array.

`get_feature_names_out(input_features=None)`[\[source\]](#)

Get output feature names for transformation.

Parameters:***input_features : array-like of str or None, default=None***

Input features.

- If `input_features` is `None`, then `feature_names_in_` is used as feature names in. If `feature_names_in_` is not defined, then the following input feature names are generated: `["x0", "x1", ..., "x(n_features_in_ - 1)"]`.
- If `input_features` is an array-like, then `input_features` must match `feature_names_in_` if `feature_names_in_` is defined.

Returns:***feature_names_out : ndarray of str objects***

Same as input features.

`get_params(deep=True)`[\[source\]](#)

Get parameters for this estimator.

Parameters:***deep : bool, default=True***If `True`, will return the parameters for this estimator and contained subobjects that are estimators.**Returns:*****params : dict***

Parameter names mapped to their values.

`inverse_transform(X, copy=None)`[\[source\]](#)

Scale back the data to the original representation.

Parameters:**X : {array-like, sparse matrix} of shape (n_samples, n_features)**

The data used to scale along the features axis.

copy : bool, default=None

Copy the input X or not.

Returns:**X_tr : {ndarray, sparse matrix} of shape (n_samples, n_features)**

Transformed array.

partial_fit(X, y=None, sample_weight=None)[\[source\]](#)

Online computation of mean and std on X for later scaling.

All of X is processed as a single batch. This is intended for cases when [fit](#) is not feasible due to very large number of n_samples or because X is read from a continuous stream.

The algorithm for incremental mean and std is given in Equation 1.5a,b in Chan, Tony F., Gene H. Golub, and Randall J. LeVeque. "Algorithms for computing the sample variance: Analysis and recommendations." The American Statistician 37.3 (1983): 242-247:

Parameters:**X : {array-like, sparse matrix} of shape (n_samples, n_features)**

The data used to compute the mean and standard deviation used for later scaling along the features axis.

y : None

Ignored.

sample_weight : array-like of shape (n_samples,), default=None

Individual weights for each sample.

New in version 0.24: parameter *sample_weight* support to StandardScaler.**Returns:****self : object**

Fitted scaler.

set_params(params)**[\[source\]](#)

Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects (such as [Pipeline](#)). The latter have parameters of the form <component>__<parameter> so that it's possible to update each component of a nested object.

Parameters:****params : dict**

Estimator parameters.

Returns:**self : estimator instance**

Estimator instance.

transform(X, copy=None)[\[source\]](#)

Perform standardization by centering and scaling.

Parameters:

X : {array-like, sparse matrix of shape (n_samples, n_features)}

The data used to scale along the features axis.

copy : bool, default=None


Copy the input X or not.

Returns:

X_tr : {ndarray, sparse matrix} of shape (n_samples, n_features)

Transformed array.


Examples using sklearn.preprocessing.StandardScaler




[Release Highlights
for scikit-learn 1.1](#)




[Release Highlights
for scikit-learn 1.0](#)



[Release Highlights
for scikit-learn 0.23](#)




[Release Highlights
for scikit-learn 0.22](#)




[Classifier
comparison](#)




[A demo of K-Means
clustering on the
handwritten digits
data](#)




[Comparing different
clustering algorithms
on toy datasets](#)




[Comparing different
hierarchical linkage
methods on toy
datasets](#)




[Demo of DBSCAN
clustering algorithm](#)




[Principal Component
Regression vs Partial
Least Squares
Regression](#)




[Factor Analysis \(with
rotation\) to visualize
patterns](#)




[Combine predictors
using stacking](#)








[Faces recognition
example using
eigenfaces and
SVMs](#)




[Prediction Latency](#)




[Comparing Linear
Bayesian Regressors](#)




[L1 Penalty and Sparsity in Logistic Regression](#)



[Lasso model selection via information criteria](#)




[Lasso model selection: AIC-BIC / cross-validation](#)




[MNIST classification using multinomial logistic + L1](#)




[Poisson regression and non-normal loss](#)




[Tweedie regression on insurance claims](#)




[Common pitfalls in the interpretation of coefficients of linear models](#)




[Advanced Plotting With Partial Dependence](#)



[Displaying Pipelines](#)




[Visualizations with Display Objects](#)




[Detection error tradeoff \(DET\) curve](#)



[Precision-Recall](#)



[Comparing Nearest Neighbors with and without Neighborhood Components](#)




[Dimensionality Reduction with Neighborhood Components Analysis](#)




[Varying regularization in Multi-layer Perceptron](#)



[Column Transformer with Mixed Types](#)



[Pipelining: chaining a PCA and a logistic regression](#)



[Compare the effect of different scalers on data with outliers](#)



[Feature discretization](#)



[Importance of Feature Scaling](#)



[RBF SVM parameters](#)



[SVM-Anova: SVM with univariate feature selection](#)

